

# Advanced Topics in Deep Reinforcement Learning

## *Exploration*



# Why Explore?



Exploration: (Strategically) deviating from the policy to improve learning later on

# Why Explore?

Exploration: (Strategically) deviating from the policy to improve learning later on

- Find important states (e.g. goals)

# Why Explore?

Exploration: (Strategically) deviating from the policy to improve learning later on

- Find important states (e.g. goals)
- Check alternative paths

# Why Explore?

Exploration: (Strategically) deviating from the policy to improve learning later on

- Find important states (e.g. goals)  
Increase state coverage
- Check alternative paths

# Why Explore?

Exploration: (Strategically) deviating from the policy to improve learning later on

- Find important states (e.g. goals)

Increase state coverage

- Check alternative paths

Avoid local minima

# Why Explore?

Exploration: (Strategically) deviating from the policy to improve learning later on

- Find important states (e.g. goals)

Increase state coverage

- Check alternative paths

Avoid local minima

=> Exploration increases data quality

# What Is High Data Quality?




- Data quality is **problem dependent**
  - State coverage is important when learning policies with many steps
  - Avoiding local minima is relevant in most reward structures
  - The degree of both varies (tradeoff with efficiency)



# What Is High Data Quality?

- Data quality is **problem dependent**
  - State coverage is important when learning policies with many steps
  - Avoiding local minima is relevant in most reward structures
  - The degree of both varies (tradeoff with efficiency)

Examples:



# Strategies for Exploration

- Randomness (e.g. action noise, e-greedy)

# Strategies for Exploration

- Randomness (e.g. action noise, e-greedy)

Easy to implement, inefficient

# Strategies for Exploration

- Randomness (e.g. action noise, e-greedy)  
Easy to implement, inefficient
- Random sequences (e.g. ez-greedy)

# Strategies for Exploration

- Randomness (e.g. action noise, e-greedy)

Easy to implement, inefficient

- Random sequences (e.g. ez-greedy)

Easy to implement, diverse sequences, inefficient, broad state coverage only

# Strategies for Exploration

- Randomness (e.g. action noise, e-greedy)

Easy to implement, inefficient

- Random sequences (e.g. [ez-greedy](#))

Easy to implement, diverse sequences, inefficient, broad state coverage only

- State novelty (e.g. [RND](#), [NovelD](#), [E3B](#), count-based methods)

# Strategies for Exploration

- Randomness (e.g. action noise, e-greedy)  
Easy to implement, inefficient
- Random sequences (e.g. [ez-greedy](#))  
Easy to implement, diverse sequences, inefficient, broad state coverage only
- State novelty (e.g. [RND](#), [NovelD](#), [E3B](#), count-based methods)  
Maximizes state coverage, hard to measure, computational overhead

# Strategies for Exploration

- Randomness (e.g. action noise, e-greedy)  
Easy to implement, inefficient
- Random sequences (e.g. [ez-greedy](#))  
Easy to implement, diverse sequences, inefficient, broad state coverage only
- State novelty (e.g. [RND](#), [NovelD](#), [E3B](#), count-based methods)  
Maximizes state coverage, hard to measure, computational overhead
- Intrinsic rewards (e.g. [ICM](#), [RIDE](#), meta-learned rewards)



# Strategies for Exploration

- Randomness (e.g. action noise, e-greedy)  
Easy to implement, inefficient
- Random sequences (e.g. [ez-greedy](#))  
Easy to implement, diverse sequences, inefficient, broad state coverage only
- State novelty (e.g. [RND](#), [NovelD](#), [E3B](#), count-based methods)  
Maximizes state coverage, hard to measure, computational overhead
- Intrinsic rewards (e.g. [ICM](#), [RIDE](#), meta-learned rewards)  
Can eliminate local minima, hard to construct, often domain specific

# Strategies for Exploration

- Randomness (e.g. action noise, e-greedy)  
Easy to implement, inefficient
- Random sequences (e.g. [ez-greedy](#))  
Easy to implement, diverse sequences, inefficient, broad state coverage only
- State novelty (e.g. [RND](#), [NovelD](#), [E3B](#), count-based methods)  
Maximizes state coverage, hard to measure, computational overhead
- Intrinsic rewards (e.g. [ICM](#), [RIDE](#), meta-learned rewards)  
Can eliminate local minima, hard to construct, often domain specific
- Model/ensemble disagreement (e.g. for [DQN](#) and [PPO](#))

# Strategies for Exploration

- Randomness (e.g. action noise,  $\epsilon$ -greedy)  
Easy to implement, inefficient
- Random sequences (e.g. [ez-greedy](#))  
Easy to implement, diverse sequences, inefficient, broad state coverage only
- State novelty (e.g. [RND](#), [NovelD](#), [E3B](#), count-based methods)  
Maximizes state coverage, hard to measure, computational overhead
- Intrinsic rewards (e.g. [ICM](#), [RIDE](#), meta-learned rewards)  
Can eliminate local minima, hard to construct, often domain specific
- Model/ensemble disagreement (e.g. for [DQN](#) and [PPO](#))  
Exploration around current policy, high computational overhead, not ideal for state coverage

# Exploration via Randomness

- Option 1: take random action with a certain probability
- Option 2: add randomly sampled noise to prediction

# Exploration via Randomness



- Option 1: take random action with a certain probability
- Option 2: add randomly sampled noise to prediction

Either way, noise should usually **decrease** during training

# Exploration via Randomness

- Option 1: take random action with a certain probability
- Option 2: add randomly sampled noise to prediction



Either way, noise should usually **decrease** during training



# Exploration via Random Sequences

Instead of sampling a step sample a **full sequence** of steps

Example assuming we sample the same action for 5 steps:



# Exploration via State Novelty

Why look for novelty?

- Exploring the whole state space
- Either focus on good performance or states we haven't seen very often



# Exploration via State Novelty

## Why look for novelty?

- Exploring the whole state space
- Either focus on good performance or states we haven't seen very often

## Why does it work?

- Avoids spending time in mediocre states
- Still helps cover the state space

# Measuring State Novelty

- Counting state visitation (limited to discrete state spaces)

# Measuring State Novelty


- Counting state visitation (limited to discrete state spaces)
- Approximating counts (in continuous state spaces)

# Measuring State Novelty

- Counting state visitation (limited to discrete state spaces)
- Approximating counts (in continuous state spaces)
- Prediction-error based novelty


# Measuring State Novelty

- Counting state visitation (limited to discrete state spaces)
- Approximating counts (in continuous state spaces)
- Prediction-error based novelty



# Exploration via State Novelty

Assume we learned to walk along the border to the goal:



# Exploration via Intrinsic Rewards

We alter the rewards instead of the policy:

- The agent will learn exploration through the rewards
- No need to choose an explicit exploration strategy
- Algorithm-agnostic
- Can encode very specific exploration behavior

# Designing Intrinsic Rewards


- Novelty-based methods often use rewards




# Designing Intrinsic Rewards

- Novelty-based methods often use rewards
- Rewards via sub-goals (e.g. [AMIGo](#))

Teacher proposes sub-goal

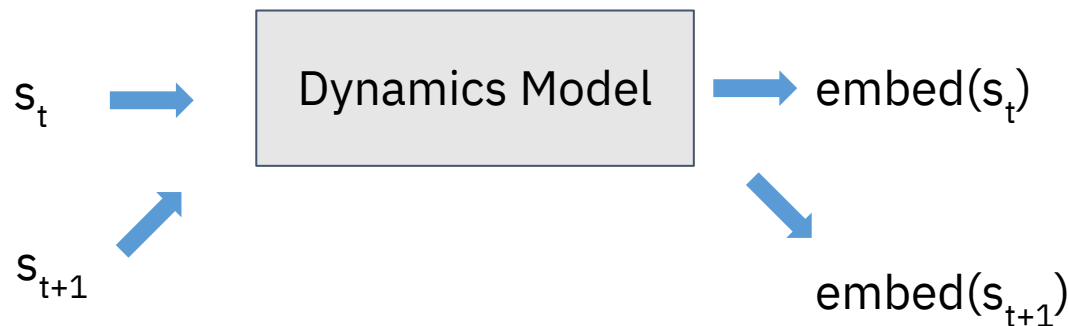


Difficulty increases over time




# Designing Intrinsic Rewards

- Novelty-based methods often use rewards
- Rewards via sub-goals (e.g. [AMIGo](#))
- Impact on learning (e.g. [ICM](#), [RIDE](#))




# Designing Intrinsic Rewards

- Novelty-based methods often use rewards
- Rewards via sub-goals (e.g. [AMIGo](#))
- Impact on learning (e.g. [ICM](#), [RIDE](#))



# Designing Intrinsic Rewards

- Novelty-based methods often use rewards
- Rewards via sub-goals (e.g. [AMIGo](#))
- Impact on learning (e.g. [ICM](#), [RIDE](#))



# Designing Intrinsic Rewards

- Novelty-based methods often use rewards
- Rewards via sub-goals (e.g. [AMIGo](#))
- Impact on learning (e.g. [ICM](#), [RIDE](#))
- Meta-learning reward functions


# Designing Intrinsic Rewards

- Novelty-based methods often use rewards
- Rewards via sub-goals (e.g. [AMIGo](#))
- Impact on learning (e.g. [ICM](#), [RIDE](#))
- Meta-learning reward functions

Designing intrinsic rewards is **very hard** and often **domain-specific**!

# Exploration via Ensembles

**Idea:** if we train multiple policies on the same data, disagreement between the predictions can be a useful signal that data is insufficient



Region of Disagreement

# My Understanding of RL Algorithms

- ❑ I understand what exploration accomplishes
- ❑ I can describe at least 3 ways of exploring
- ❑ I could implement at least one exploration algorithm
- ❑ I can describe each exploration idea
- ❑ I know the pros and cons of different exploration methods
- ❑ I can describe one algorithm per idea
- ❑ I can compare exploration algorithms
- ❑ I can explain which idea would work for a given problem

