

Online Meta-Exploration from Offline Data in Realistic Robotic Tasks

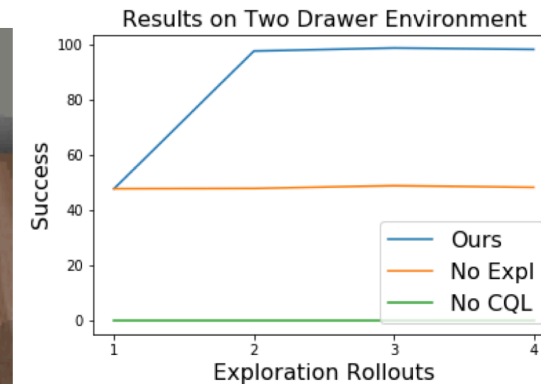

[Rafailov et. al. 2021]



Problem and Approach [Rafailov et. al. 2021]

- **Challenges from Real-World RL:**
 - High sample complexity
 - Distribution shifts
 - High-dimensional inputs (e.g., images)
 - Spurious memorization instead of learning
- **Current Solutions and Limitations:**
 - Meta-learning aids agents for quick adaptation to new tasks
 - requires a lot of data and is hard to optimize in high-dimensional spaces
- **Proposed Approach:**
 - Operates from images with sparse rewards
 - Meta-exploration from offline-data
 - Using variational filtering for representation learning (=)
 - Latent conservative model-free policy optimization (=)

Methodology and Results [Rafailov et. al. 2021]



- **Synthetic Exploration**
 - Creating extended exploration rollouts by mixing trajectories from different tasks
 - Assigning target tasks & zeroing out rewards for non-target tasks (to avoid memorization/correlation)
- **Variational Model Training**
 - Training a high-dimensional variational model to infer environment and tasks states from image sequences
- **Latent Conservative Model-Free policy Optimization**
 - Using CQL (Conservative Q-Learning) to slow over-estimation bias and distribution shift issues (primarily from offline RL)
 - **CQL** = Penalizes Q-Values of actions which are not well-supported by the offline data, causes reliable policy learning
- **Results**
 - Successfully solved a meta-learning task from a realistic problem
 - Outperformed naive combinations of previous approaches which already failed on the same task

Sources

[Rafailov et. al. 2021] <https://offline-rl-neurips.github.io/2021/pdf/54.pdf>