

Proximal Curriculum Learning with State Novelty

- ▶ Use zone of proximal development (pick start state with probability of success $PoS \approx 0.5$) [Tzannetos et al. 2023]
- ▶ approximate PoS via value-function [Tzannetos et al. 2023]

Proximal Curriculum Learning with State Novelty

- ▶ Use zone of proximal development (pick start state with probability of success $PoS \approx 0.5$) [Tzannetos et al. 2023]
 - ▶ approximate PoS via value-function [Tzannetos et al. 2023]
 - ▶ *Own Idea*: already using state to guide student, why not pick novel states
- should help value-function converge
- should still pick mostly promising states

Proximal Curriculum Learning with State Novelty

- ▶ Use zone of proximal development (pick start state with probability of success $PoS \approx 0.5$) [Tzannetos et al. 2023]
- ▶ approximate PoS via value-function [Tzannetos et al. 2023]
- ▶ *Own Idea*: already using state to guide student, why not pick novel states
 - should help value-function converge
 - should still pick mostly promising states
- ▶ trade off via hyperparameter between PoS and novelty
- ▶ alternatively incorporate state novelty as intrinsic reward: directly influences PoS

Practical Considerations

- ▶ *Implementation Overhad*: no given code, but modifications (including hyperparameters) to stable-baselines listed for [Tzannetos et al. 2023]

Practical Considerations

- ▶ *Implementation Overhad*: no given code, but modifications (including hyperparameters) to stable-baselines listed for [Tzannetos et al. 2023]
- ▶ *Compute Overhead*: should be reasonably cheap depending on the state novelty algorithm and environment

Practical Considerations

- ▶ *Implementation Overhad*: no given code, but modifications (including hyperparameters) to stable-baselines listed for [Tzannetos et al. 2023]
- ▶ *Compute Overhead*: should be reasonably cheap depending on the state novelty algorithm and environment
- ▶ *Reproducibility*: Should reproduce given limited modifications to stable-baselines. If not try to use rollouts to approximate *PoS*.

Practical Considerations

- ▶ *Implementation Overhad*: no given code, but modifications (including hyperparameters) to stable-baselines listed for [Tzannetos et al. 2023]
- ▶ *Compute Overhead*: should be reasonably cheap depending on the state novelty algorithm and environment
- ▶ *Reproducibility*: Should reproduce given limited modifications to stable-baselines. If not try to use rollouts to approximate *PoS*.
- ▶ MiniGrid should be good, sparse and many states